

Practical 1 :

Breadth First search Program Code :

```
#include<iostream>

#include<stdlib.h>

#include<queue> using
namespace std;

class node
{
public:
    node *left, *right;    int
    data;
};

class Breadthfs
{
public:
    node *insert(node *, int); void
    bfs(node *);
};

node *insert(node *root, int data)
// inserts a node in tree
{ if(!root)
{
    root=new node;
    root->left=NULL;
    root->right=NULL;
    root->data=data;
    return root;
}
```

```

        queue<node *> q;
q.push(root);
    While(!q.empty())
    {
        node *temp=q.front();
            q.pop();
            if(temp->left==NULL)
            {
temp->left=new node;

temp->left->left=NULL;

temp->left->right=NULL;

temp->left->data=data;

                return root;
            }
            else
            {
                q.push(temp->left);
            }
            if(temp->right==NULL)
            {

temp->right=new node;

temp->right->left=NULL;

temp->right->right=NULL;

temp->right->data=data;

return root;

            }

```

```

        else
        {
            q.push(temp->right);
        }
    }
}

void bfs(node *head)
{
    queue<node*> q;
    q.push(head);

    int qSize;
    while (!q.empty())
    {
        qSize = q.size();

#pragma omp parallel for
        //creates parallel threads
        for (int i = 0; i < qSize; i++)
        {
            node* currNode;

#pragma omp critical
            {
                currNode = q.front();
                q.pop();

                cout<<"\t"<<currNode->data;

            }

            }// prints parent node

#pragma omp critical
            {
                if(currNode->left)// push parent's left node in queue

```

```

        q.push(currNode->left);

        if(currNode->right)

            q.push(currNode->right);
        }
    }
}

int main(){
    node *root=NULL;

    int data;    char ans;

    do
    {
        cout<<"\n enter data=>";

        cin>>data;

        root=insert(root,data);

        cout<<"do you want insert one more node?";

        cin>>ans;

    }while(ans=='y' || ans=='Y');

    bfs(root);

    return 0;
}

```

Output :

```
C:\Users\pc\AppData\Local\Ti X + v

enter data=>7
do you want insert one more node?y

enter data=>4
do you want insert one more node?y

enter data=>6
do you want insert one more node?y

enter data=>1
do you want insert one more node?y

enter data=>8
do you want insert one more node?y

enter data=>9
do you want insert one more node?y

enter data=>2
do you want insert one more node?n
      7      4      6      1      8      9      2
-----
Process exited after 16.2 seconds with return value 0
Press any key to continue . . . |
```

Program For Depth First Search :

```
#include <iostream>

#include <vector>

#include <stack>

#include <omp.h>

using namespace std;

const int MAX = 100000;

vector<int> graph[MAX]; bool
visited[MAX];

void dfs(int node) {
    stack<int> s;
    s.push(node);

    while (!s.empty()) {    int
curr_node = s.top();

        if (!visited[curr_node]) {
visited[curr_node] = true;
            s.pop();
cout<<curr_node<<" ";

            #pragma omp parallel for

for (int i = 0; i < graph[curr_node].size(); i++) {
int adj_node = graph[curr_node][i];
```

```

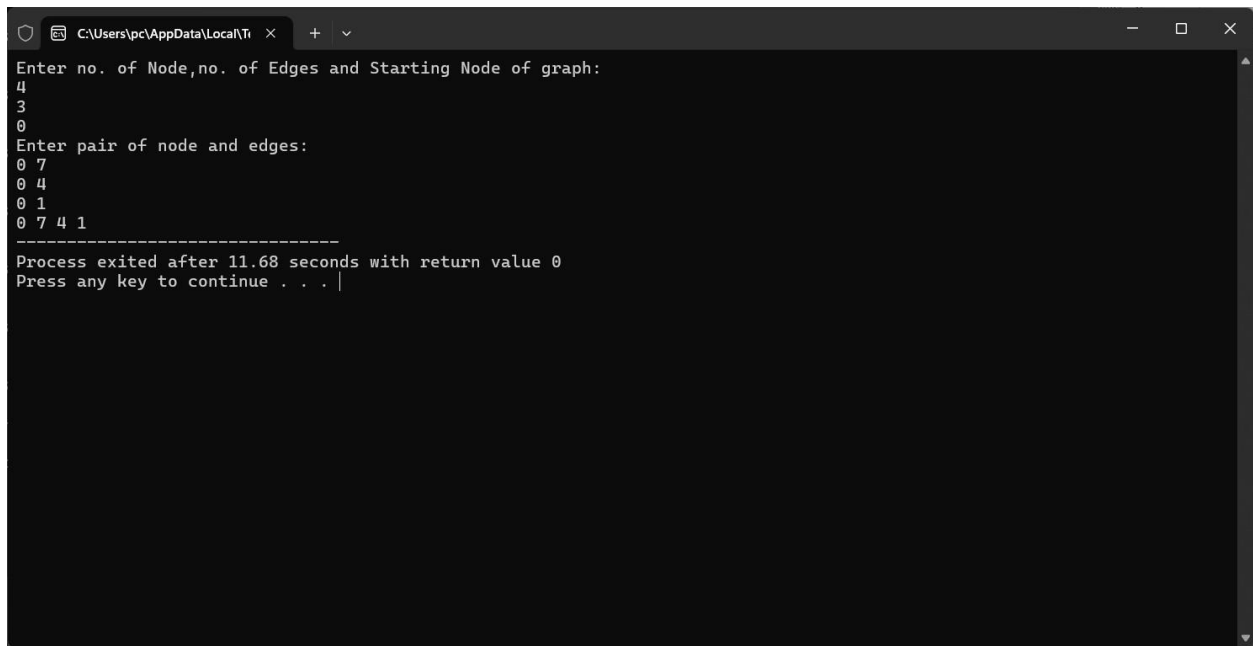
        if (!visited[adj_node]) {
s.push(adj_node);
        }    }    }    }}
int main() {
    int n, m, start_node;
    cout<<"Enter no. of Node,no. of Edges and Starting Node of graph:\n";
    cin >> n >> m >> start_node;
    //n: node,m:edges    cout<<"Enter pair
of node and edges:\n";

    for (int i = 0; i < m; i++) {
int u, v;    cin >> u >> v;

//u and v: Pair of edges
graph[u].push_back(v);
graph[v].push_back(u);
    }
    #pragma omp parallel for
    for (int i = 0; i < n; i++) {
visited[i] = false;
    }
    dfs(start_node);
    return 0;
}

```

Output :-



```
C:\Users\pc\AppData\Local\Ti x + v
Enter no. of Node,no. of Edges and Starting Node of graph:
4
3
0
Enter pair of node and edges:
0 7
0 4
0 1
0 7 4 1
-----
Process exited after 11.68 seconds with return value 0
Press any key to continue . . . |
```


Practical 2 :

Program Code For Bubble Sort :

```
#include<iostream>
```

```
#include<stdlib.h>
```

```
#include<omp.h> using
```

```
namespace std;
```

```
void bubble(int *, int); void
```

```
swap(int &, int &); void
```

```
bubble(int *a, int n)
```

```
{
```

```
    int swapped;
```

```
    for( int i = 0; i < n; i++ )
```

```
    {
```

```
        int first = i % 2;
```

```
swapped=0;
```

```
    #pragma omp parallel for shared(a,first)
```

```
    for( int j = first; j < n-1; j += 2 )
```

```
    {
```

```
        if( a[ j ] > a[ j+1 ] )
```

```
        {
```

```
            swap( a[ j ], a[ j+1 ] );
```

```
swapped=1;
```

```
        }
```

```
    }
```

```
    if(swapped==0)
```

```
break;
```

```
    }
```

```
}
```

```

void swap(int &a, int &b)
{
    int test;
    test=a;  a=b;
    b=test;
}

int main() {
    int *a,n;

    cout<<"\n enter total no of elements=>";
    cin>>n;  a=new int[n];  cout<<"\n enter
elements=>";  for(int i=0;i<n;i++)
    {
        cin>>a[i];
    }

    double start_time = omp_get_wtime(); // start timer for sequential algorithm
    bubble(a,n);  double end_time = omp_get_wtime(); // end timer for sequential
algorithm

    cout<<"\n sorted array is=>";
    for(int i=0;i<n;i++)
    {
        cout<<a[i]<<endl;
    }

    cout << "Time taken by sequential algorithm: " << end_time - start_time << " seconds" << endl;
    start_time = omp_get_wtime(); // start timer for parallel algorithm    bubble(a,n);    end_time =
omp_get_wtime(); // end timer for parallel algorithm

```

```
    cout<<"\n sorted array is=>";  
for(int i=0;i<n;i++)  
{  
    cout<<a[i]<<endl;  
}  
  
cout << "Time taken by parallel algorithm: " << end_time - start_time << " seconds" << endl;  
  
return 0;  
}
```

Output For Bubble Sort :

```
enter total no of elements=>10  
enter elements=>12  
24  
51  
23  
54  
75  
84  
56  
12  
42
```

```
sorted array is=>12
12
23
24
42
51
54
56
75
84
Time taken by sequential algorithm: 0.0149999 seconds

sorted array is=>12
12
23
24
42
51
54
56
75
84
Time taken by parallel algorithm: 0 seconds

-----
Process exited after 10.57 seconds with return value 0
Press any key to continue . . . |
```

Program for Merge Sort :

```
#include<iostream>
```

```
#include<stdlib.h>
```

```
#include<omp.h> using
```

```
namespace std;
```

```
void mergesort(int a[],int i,int j); void
```

```
merge(int a[],int i1,int j1,int i2,int j2);
```

```
void mergesort(int a[],int i,int j)
```

```
{
```

```
    int mid;
```

```
    if(i<j)
```

```
    {
```

```
        mid=(i+j)/2;
```

```

#pragma omp parallel sections
{
    #pragma omp section
    {
        mergesort(a,i,mid);
    }
    #pragma omp section
    {
        mergesort(a,mid+1,j);
    }
}
merge(a,i,mid,mid+1,j);
}
}

void merge(int a[],int i1,int j1,int i2,int j2)
{
    int temp[1000];
    int i,j,k;   i=i1;
j=i2;    k=0;
while(i<=j1 && j<=j2)
{
    if(a[i]<a[j])
    {
        temp[k++]=a[i++];
    }
    else
    {
        temp[k++]=a[j++];
    }
}
}

```

```

        }
    }
    while(i<=j1)
    {
        temp[k++]=a[i++];
    }
    while(j<=j2)
    {
        temp[k++]=a[j++];
    }
    for(i=i1,j=0;i<=j2;i++,j++)
    {
        a[i]=temp[j];
    }
}

int main()
{
    int
    *a,n,i;

    double start_time, end_time, seq_time, par_time;
    cout<<"\n enter total no of elements=>";    cin>>n;
    a= new int[n];    cout<<"\n enter elements=>";
    for(i=0;i<n;i++)
    {
        cin>>a[i];
    }

    // Sequential algorithm    start_time =
    omp_get_wtime();    mergesort(a, 0, n-1);    end_time
    = omp_get_wtime();    seq_time = end_time -

```

```

start_time;  cout << "\nSequential Time: " <<
seq_time << endl;

// Parallel algorithm
start_time = omp_get_wtime();
#pragma omp parallel
{
    #pragma omp single
    {
        mergesort(a, 0, n-1);
    }
}

end_time = omp_get_wtime();  par_time =
end_time - start_time;  cout << "\nParallel Time:
" << par_time << endl;

cout<<"\n sorted array is=>";
for(i=0;i<n;i++)
{
    cout<<"\n"<<a[i];
}

return 0;
}

Output for Merge Sort :
```

```
enter total no of elements=>15
enter elements=>21
52
58
564
48456
6548
1234
457
12
548
456
774
548
654
245

Sequential Time: 0
Parallel Time: 0

sorted array is=>
12
21
52
58
245
456
457
548
548
564
654
774
1234
6548
48456
-----
Process exited after 22.38 seconds with return value 0
```

Practical 3 :

Program for Parallel Reduction :

```
#include <iostream>
```

```
#include <vector>
```

```
#include <omp.h>
```

```
#include <climits>
```

```
using namespace std;
```

```
void min_reduction(vector<int>& arr) { int  
min_value = INT_MAX;  
    #pragma omp parallel for reduction(min: min_value)  
    for (int i = 0; i < arr.size(); i++) {  
if (arr[i] < min_value) {  
min_value = arr[i];  
    }  
    }  
    cout << "Minimum value: " << min_value << endl;  
}
```

```
void max_reduction(vector<int>& arr) { int  
max_value = INT_MIN;  
    #pragma omp parallel for reduction(max: max_value)  
    for (int i = 0; i < arr.size(); i++) {  
if (arr[i] > max_value) {  
max_value = arr[i];  
    }  
    }  
    cout << "Maximum value: " << max_value << endl;
```

```
}
```

```
void sum_reduction(vector<int>& arr) { int  
sum = 0;  
    #pragma omp parallel for reduction(+: sum)  
    for (int i = 0; i < arr.size(); i++) {  
sum += arr[i];  
    }  
    cout << "Sum: " << sum << endl;  
}
```

```
void average_reduction(vector<int>& arr) { int  
sum = 0;  
    #pragma omp parallel for reduction(+: sum)  
    for (int i = 0; i < arr.size(); i++) {  
sum += arr[i];  
    }  
    cout << "Average: " << (double)sum / arr.size() << endl;  
}
```

```
int main() {  
vector<int> arr;  
arr.push_back(5);  
arr.push_back(2);  
arr.push_back(9);  
arr.push_back(1);  
arr.push_back(7);  
arr.push_back(6);  
arr.push_back(8);
```

```
arr.push_back(3);
```

```
arr.push_back(4);
```

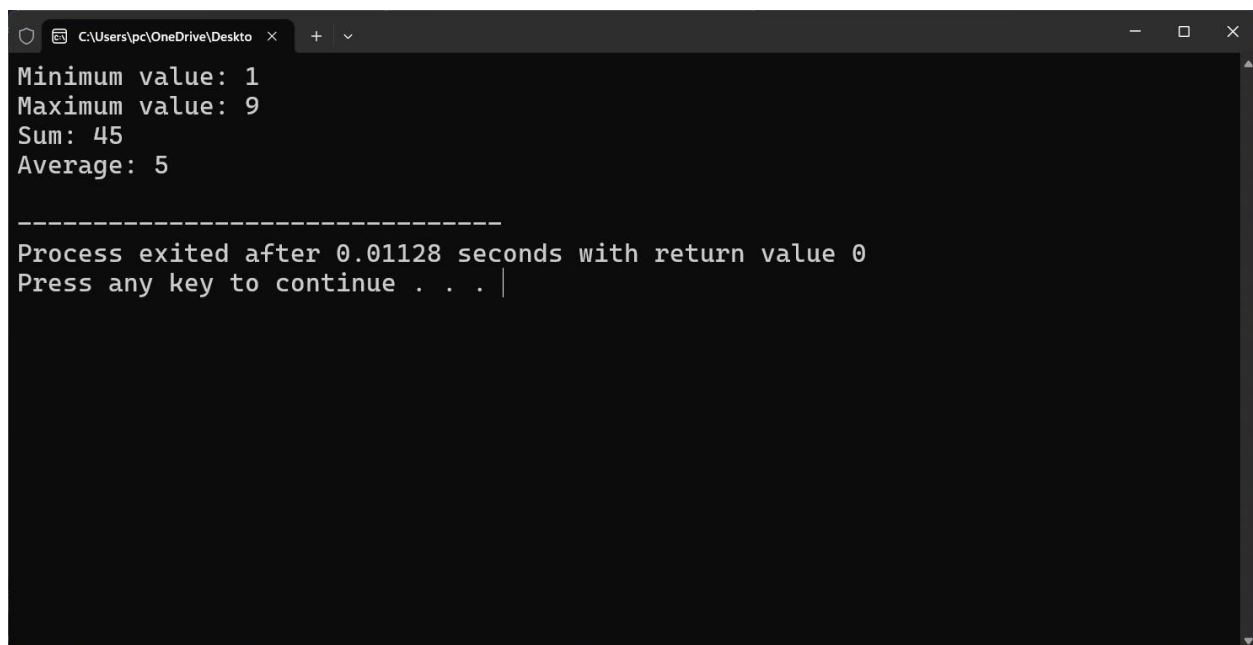
```
    min_reduction(arr); max_reduction(arr);
```

```
    sum_reduction(arr);
```

```
    average_reduction(arr);
```

```
}
```

Output :



```
Minimum value: 1
Maximum value: 9
Sum: 45
Average: 5

-----
Process exited after 0.01128 seconds with return value 0
Press any key to continue . . .
```

Practical 04 :

Program for Matrix multiplication using CUDA :

```
#include <cuda_runtime.h>
#include <iostream>

__global__ void matmul(int* A, int* B, int* C, int N) {
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
    int Col = blockIdx.x * blockDim.x + threadIdx.x;    if
    (Row < N && Col < N) {        int Pvalue = 0;
        for (int k = 0; k < N; k++) {
            Pvalue += A[Row * N + k] * B[k * N + Col];
        }
        C[Row * N + Col] = Pvalue;
    }
}

int main() {
    int N = 128;
    int size = N * N * sizeof(int);
    int* A, * B, * C;    int* dev_A, *
    dev_B, * dev_C;
    cudaMallocHost(&A, size);
    cudaMallocHost(&B, size);
    cudaMallocHost(&C, size);
    cudaMalloc(&dev_A, size);
    cudaMalloc(&dev_B, size);
    cudaMalloc(&dev_C, size);

    // Initialize matrices A and B
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            A[i * N + j] = i * N + j;
            B[i * N + j] = j * N + i;
        }
    }

    cudaMemcpy(dev_A, A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(dev_B, B, size, cudaMemcpyHostToDevice);

    dim3 dimBlock(16, 16);
    dim3 dimGrid(N / dimBlock.x, N / dimBlock.y);

    matmul << <dimGrid, dimBlock >> > (dev_A, dev_B, dev_C, N);

    cudaMemcpy(C, dev_C, size, cudaMemcpyDeviceToHost);

    // Print the result    for (int i = 0; i
    < 10; i++) {        for (int j = 0; j < 10;
    j++) {            std::cout << C[i * N + j]
    << " ";
        }
        std::cout << std::endl;
    }
}
```

```

    // Free memory
    cudaFree(dev_A);
    cudaFree(dev_B);
    cudaFree(dev_C);
    cudaFreeHost(A);
    cudaFreeHost(B);
    cudaFreeHost(C);

    return 0;
}

```

Output :

```

Microsoft Visual Studio Debug Console
690880 1731264 2771648 3812032 4852416 5892800 6933184 7973568 9013952 10054336
1731264 4868800 8006336 11143872 14281408 17418944 20556480 23694016 26831552 29969088
2771648 8006336 13241024 18475712 23710400 28945088 34179776 39414464 44649152 49883840
3812032 11143872 18475712 25807552 33139392 40471232 47803072 55134912 62466752 69798592
4852416 14281408 23710400 33139392 42568384 51997376 61426368 70855360 80284352 89713344
5892800 17418944 28945088 40471232 51997376 63523520 75049664 86575808 98101952 109628096
6933184 20556480 34179776 47803072 61426368 75049664 88672960 102296256 115919552 129542848
7973568 23694016 39414464 55134912 70855360 86575808 102296256 118016704 133737152 149457600
9013952 26831552 44649152 62466752 80284352 98101952 115919552 133737152 151554752 169372352
10054336 29969088 49883840 69798592 89713344 109628096 129542848 149457600 169372352 189287104

C:\Users\pc\source\repos\CudaRuntime1\x64\Debug\CudaRuntime1.exe (process 18736) exited with code 0.
To automatically close the console when debugging stops, enable Tools->Options->Debugging->Automatically close the console when debugging stops.
Press any key to close this window . . .|

```

Program for Vector Addition using CUDA :

```
#include <iostream>
#include <cuda_runtime.h>

using namespace std;

__global__ void addVectors(int* A, int* B, int* C, int n)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < n)
    {
        C[i] = A[i] + B[i];
    }
}

int main() {
    int n = 1000000;    int*
    A, * B, * C;    int size = n *
    sizeof(int);

    // Allocate memory on the host
    cudaMallocHost(&A, size);    cudaMallocHost(&B,
    size);    cudaMallocHost(&C, size);

    // Initialize the vectors
    for (int i = 0; i < n; i++)
    {
        A[i] = i;
        B[i] = i * 2;
    }

    // Allocate memory on the device
    int* dev_A, * dev_B, * dev_C;
    cudaMalloc(&dev_A, size);
    cudaMalloc(&dev_B, size);
    cudaMalloc(&dev_C, size);

    // Copy data from host to device
    cudaMemcpy(dev_A, A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(dev_B, B, size, cudaMemcpyHostToDevice);

    // Launch the kernel
    int blockSize = 256;
    int numBlocks = (n + blockSize - 1) / blockSize;
    addVectors << <numBlocks, blockSize >> > (dev_A, dev_B, dev_C, n);

    // Copy data from device to host
    cudaMemcpy(C, dev_C, size, cudaMemcpyDeviceToHost);

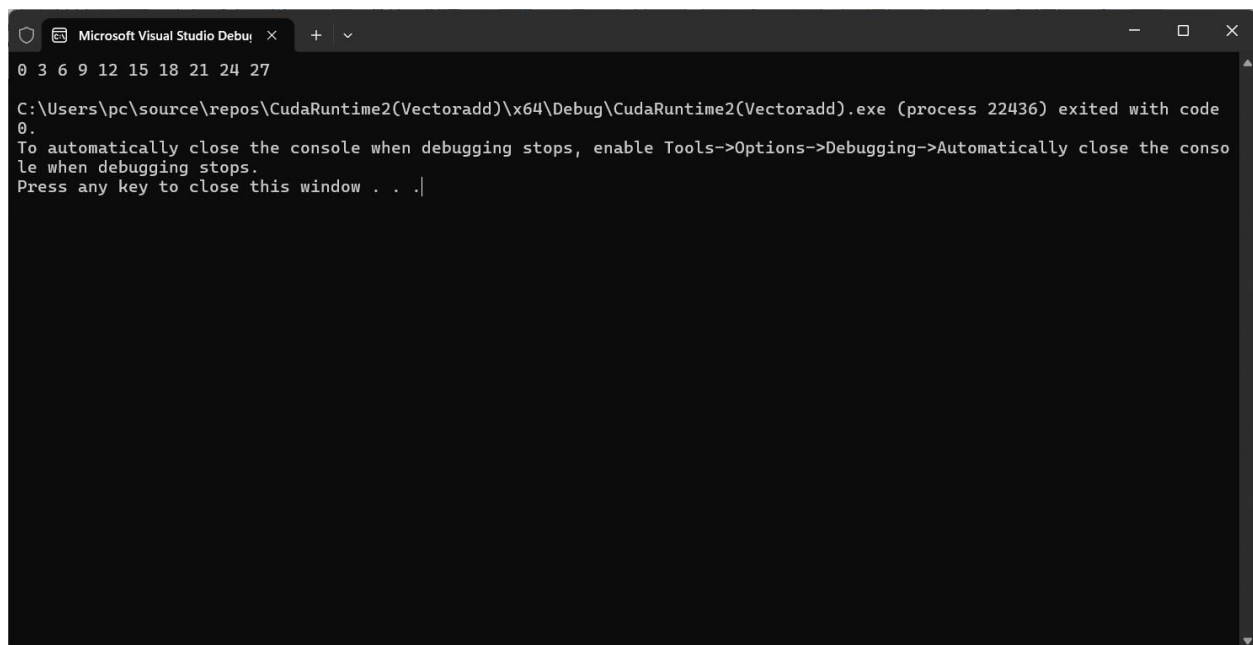
    // Print the results    for
    (int i = 0; i < 10; i++)
    {
        cout << C[i] << " ";
    }
}
```

```
    cout << endl;

    // Free memory
    cudaFree(dev_A);    cudaFree(dev_B);
    cudaFree(dev_C);    cudaFreeHost(A);
    cudaFreeHost(B);    cudaFreeHost(C);

    return 0;
}
```

Output :



```
Microsoft Visual Studio Debug Console
0 3 6 9 12 15 18 21 24 27
C:\Users\pc\source\repos\CudaRuntime2(Vectoradd)\x64\Debug\CudaRuntime2(Vectoradd).exe (process 22436) exited with code 0.
To automatically close the console when debugging stops, enable Tools->Options->Debugging->Automatically close the console when debugging stops.
Press any key to close this window . . .|
```